

School of Computing and Information Systems  
**COMP30023: Computer Systems**

Tutorial Week 10

**Process communication, scheduling and memory management**

1. What is a race condition? What is a deadlock?
2. In the lectures we saw that the priority inversion problem can happen with processes. Can the priority inversion problem happen with user-level threads? Why or why not?
3. In lectures we saw a concept of a resource graph and a circular chain that leads to a deadlock. Give an example to show that the set of processes deadlocked can include processes that are not in a circular chain in the corresponding resource allocation graph.
4. Can a measure of whether a process is likely to be CPU bound or I/O bound be determined by analysing source code? How can this be determined at run time?
5. Five batch jobs, A through E, arrive at a computer centre at almost the same time. They have estimated running times of 10, 6, 2, 4, and 8 minutes. Their (externally determined) priorities are 3, 5, 2, 1, and 4, respectively, with 5 being the highest priority. For each of the following scheduling algorithms, determine the mean process turnaround time. Ignore process switching overhead. (a) Round robin with quantum of 6 minutes. (b) Priority scheduling. (c) First-come, first-served (run in order 10, 6, 2, 4, 8). (d) Shortest job first. For (b) through (d), assume that only one job at a time runs, until it finishes. All jobs are completely CPU bound.

6. Give an arithmetic expression connecting a given virtual address with the corresponding page number and offset. What does this tell you about page sizes that are not powers of two?
7. How much physical memory do you need to store a page table describing (a) 512 KB (b) 4 GB given a page size of 512 bytes? How much do you need given a page size of 8 Kb? Assume that Page Table Entries (PTEs) are 4 bytes in either case.
8. You are given the following data about a virtual memory system:
  - (a) The TLB can hold 1024 entries and can be accessed in 1 clock cycle (1 nsec).
  - (b) A page table entry can be found in 100 clock cycles or 100 nsec.
  - (c) The average page replacement time is 6 msec.

If page references are handled by the TLB 99% of the time, and only 0.01% lead to a page fault, what is the effective address-translation time?

9. In the lectures we saw that there are several strategies for choosing which free memory block (hole) to assign to a process. First/best/worst fit algorithms operate as follows: given a memory request of size  $X$  choose the first/smallest/largest hole that fits  $X$ , correspondingly.

Consider a swapping system in which memory consists of the following hole sizes in memory order: 10 MB, 4 MB, 20 MB, 18 MB, 7 MB, 9 MB, 12 MB, and 15 MB. Which hole is taken for successive segment requests of (a) 12 MB (b) 10 MB (c) 9 MB for first fit? Now repeat the question for best fit, worst fit.

10. Working set page replacement algorithm will first try to evict pages that are not in the working set before evicting pages in the working set. Under what type of access locality does a working set algorithm perform best (i.e., minimise number of page faults) and why?

### Additional questions

11. Recall the multi-threaded web-server scenario during Week 9. Another design could be formulated via consumer-producer problem as follows. There is a dispatcher and a worker thread that both have access to a buffer of requests. The dispatcher writes a request to a buffer and the worker thread reads the request, removes it and serves it by returning the corresponding page. If the buffer is empty, the worker sleeps until the

dispatcher wakes it up. If the buffer is full, the dispatcher sleeps until worker wakes it up. See the pseudo-code below.

Under which circumstances this code can lead to worker and dispatcher sleeping at the same time?

Can you think of a way to resolve the problem?

```
#define N 100 /* number of slots in the buffer */
int count = 0; /* number of requests in the buffer */

void dispatcher(void) {
    int request;
    while (TRUE) { /* repeat forever */
        request = add_request( ); /* add next request */
        if (count == N) sleep(); /* if buffer is full, go to sleep */
        insert_request(request); /* put request in buffer */
        count = count + 1; /* increment counter */
        if (count == 1) wakeup(worker); /* was buffer empty? */
    }
}

void worker(void) {
    int request;
    while (TRUE) { /* repeat forever */
        if (count == 0) sleep(); /* if buffer is empty, go to sleep */
        request = remove_request(); /* take request out of buffer */
        count = count - 1; /* decrease the count */
        if (count == N - 1) wakeup(dispatcher); /* was buffer full? */
        serve_request(request);
    }
}
```

12. Consider a system with 32-bit virtual (logical) addresses, 24-bit physical addresses, and the system uses paging to manage memory. A page frame holds 4096 bytes. What is the maximum number of entries in the page table?
13. What is meant by the term *thrashing*?